

A PROJECT REPORT

on

TIME SERIES MACHINE LEARNING

Submitted to

KIIT DEEMED TO BE UNIVERSITY

In Partial Fulfilment of the Requirement for the Award of

BACHELOR'S DEGREE IN
COMPUTER SCIENCE & ENGINEERING

BY

ADARSH NARAYAN	22052352
AKASH KUMAR	22051659
AMAN PATHAK	22051662
AMIT KUMAR	22053747

UNDER THE GUIDANCE OF

Prof. Jyotiprakash MISHRA



SCHOOL OF COMPUTER ENGINEERING
KALINGA INSTITUTE OF INDUSTRIAL
TECHNOLOGY

BHUBANESWAR, ODISHA - 751024

April 2025

Acknowledgements

We would like to express our sincere gratitude to our guide, Prof. Jyotiprakash Mishra, for his invaluable guidance, support, and encouragement throughout this project. We also extend our thanks to the faculty members of the School of Computer Engineering, KIIT Deemed to be University, for their insightful lectures and constant motivation. Lastly, we are grateful to our families and friends for their continuous support and encouragement.

Adarsh Narayan

Akash Kumar

Aman Pathak

Amit Kumar

Abstract

Time series forecasting plays a crucial role in various domains, enabling accurate predictions based on historical patterns. This project explores different time series machine learning techniques applied to four diverse datasets, each addressing a unique problem statement. The first dataset, Household Power Consumption, involves forecasting hourly and daily energy usage patterns to enhance energy efficiency. The second dataset, Predictive Maintenance, focuses on detecting potential machine failures using sensor data, contributing to proactive maintenance strategies. The third dataset, Walmart and Rossmann Store Sales, aims to predict retail sales while accounting for seasonal trends, promotional effects, and holidays. Lastly, the Historical Stock Market Data dataset is used to predict closing prices and identify potential trading signals, assisting in financial decision-making.

The project follows a structured approach, beginning with data wrangling techniques such as timestamp conversion, handling missing data, and resampling to maintain consistency. Feature engineering is employed to enhance predictive performance, incorporating moving averages, rolling windows, and external data sources like holidays, weather conditions, and macroeconomic indicators. A comprehensive model comparison is conducted, evaluating traditional statistical models (ARIMA/SARIMA), machine learning methods (Random Forest, XGBoost), and deep learning architectures (LSTM, GRU). Model performance is assessed using error metrics such as Mean Squared Error (MSE), Root Mean Squared Error (RMSE), Mean Absolute Error (MAE), and Mean Absolute Percentage Error (MAPE). Additionally, explainability techniques such as SHAP values and feature importance plots are applied to ensemble models, providing insights into the key drivers of predictions. Visualizations, including forecasted vs. actual values, seasonal trends, and anomaly detection, further illustrate the effectiveness of the chosen models. The results highlight the strengths and limitations of various approaches, demonstrating the applicability of machine learning and deep learning models in real-world time series forecasting problems.

Contents

1	Time Series Machine Learning for Household Power Consumption	6
1.1	Introduction to Machine Learning	6
1.1.1	Definitions of Machine Learning	6
1.1.2	Sub-Categories of Machine Learning	7
1.1.3	Related Fields	7
1.1.4	Applications of Machine Learning	8
1.1.5	Relevance to Time Series Forecasting	8
1.2	Differences Between Time Series and Machine Learning	8
1.3	Autoregression	9
1.4	Sliding Window	9
1.4.1	Predicting Target or Differences	10
1.4.2	Intuition Behind First Differences	10
1.5	Data Preparation and Analysis	11
1.5.1	Cleaning the Dataset	11
1.5.2	Exploratory Data Analysis	12
1.5.3	Resampling by Time Unit	12
1.5.4	Converting to a Supervised Learning Problem	14
1.6	Statistical Models	15
1.6.1	Data Preparation and Model Configuration	15
1.6.2	Forecasting and Evaluation Metrics	16
1.6.3	Visualization of Forecasts	16
1.6.4	Discussion	17
1.7	Machine Learning Models	18
1.8	Deep Learning Models	20
1.8.1	Model Architectures	21
1.8.2	Methodology	22
1.8.3	Performance Evaluation	22
1.8.4	Discussion	24
1.9	Interpretation of the Results	26
1.9.1	Subset Dataset Analysis	26
1.9.2	Full Dataset Analysis	27
1.9.3	Comparative Insights and Practical Implications	27
2	Stock Price Prediction Analysis	30
2.1	Introduction	30
2.2	Dataset Overview	30
2.3	Models used for training	30
2.3.1	LSTM Model	30

2.3.2	Training process	30
2.3.3	Random Forest Model	31
2.3.4	SARIMA Model	31
2.4	Model evaluation	31
2.4.1	Actual vs Predicted Closing Prices	31
2.4.2	Distribution of Closing Price Differences	31
2.4.3	Model Comparison	32
2.5	Conclusion	33
3	Walmart Recruiting - Store Sales Forecasting	35
3.1	Introduction	35
3.2	Tools Used in Sales Forecasting	35
3.3	Data Overview	36
3.4	Output Analysis	37
3.5	Detailed Explanation of Tools and Outputs	39
3.5.1	ARIMA Model	39
3.5.2	XGBoost Model	39
3.5.3	LSTM Model	39
3.6	Humanized Perspective	40
3.7	Discussion	40

Chapter 1

Time Series Machine Learning for Household Power Consumption

1.1 Introduction to Machine Learning

Machine Learning (ML) is a subset of artificial intelligence that enables systems to learn from data and improve over time without explicit programming. In this chapter, we explore ML techniques applied to time series forecasting, specifically for predicting household power consumption. This section delves into the definitions, sub-categories, related fields, and practical applications of ML, setting the stage for its application to our problem.

1.1.1 Definitions of Machine Learning

Machine learning has been defined by multiple scholars in related ways, emphasizing its reliance on data and experience. Tom Mitchell, in his seminal work *Machine Learning* [1], provides a formal definition:

A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P , if its performance at tasks in T , as measured by P , improves with experience E .

This definition encapsulates algorithms like gradient descent and Q-learning, where performance improves through iterative exposure to data.

Kevin Murphy, in *Machine Learning: A Probabilistic Perspective* [2], describes ML as:

A set of methods that can automatically detect patterns in data, and then use the uncovered patterns to predict future data, or to perform other kinds of decision making under uncertainty (such as planning how to collect more data!).

This highlights ML's predictive and decision-making capabilities.

Shalev-Shwartz and Ben-David offer a simpler perspective [3]:

The term machine learning refers to the automated detection of meaningful patterns in data.

Across these definitions, the core concept is data-driven pattern recognition to solve tasks. However, the term “automatically” can mislead; ML algorithms typically require human-defined performance measures and parameters, indicating they are not fully autonomous.

As a field, ML can be defined as the study and application of such algorithms, aligning with Mitchell’s framework.

1.1.2 Sub-Categories of Machine Learning

Machine learning is often divided into three primary sub-categories, as noted by Murphy and others [2]:

- **Supervised Learning (Predictive):** The goal is to learn a mapping from inputs x to outputs y , given labeled input-output pairs. For our household power consumption problem, this involves predicting energy usage from historical data.
- **Unsupervised Learning (Descriptive):** This seeks to find interesting patterns in data without predefined labels, such as clustering consumption trends.
- **Reinforcement Learning:** Useful for learning optimal actions based on occasional reward or punishment signals, less relevant here but notable in broader ML contexts.

Additional sub-categories or taxonomies include:

- **Deep Learning:** Utilizes neural networks and algorithms like gradient descent to approximate complex functions, critical for later sections on RNNs and LSTMs.
- **Probabilistic Machine Learning:** Offers uncertainty estimation, enhancing forecast reliability.
- **Weakly Supervised Learning:** Addresses imperfect labels, a potential challenge in noisy datasets.
- **Online Learning:** Processes one data point at a time, contrasting with batch learning from datasets.

These can intersect, e.g., deep learning can be online or offline, offering flexibility in application.

1.1.3 Related Fields

A closely related field is *computational (or statistical) learning theory*, which examines the theoretical underpinnings of learning from computational and statistical perspectives. It addresses questions like, “How many samples are needed to approximate a function with a given error?”

ML also overlaps with statistics, adopting concepts like regression and hypothesis testing. While some quip that ML is “glorified statistics,” the fields differ in focus: statistics emphasizes inference and modeling, while ML prioritizes prediction and automation. For a deeper comparison, see [4].

1.1.4 Applications of Machine Learning

ML excels at automating tasks involving data and pattern recognition, historically human-dominated, such as language translation or image classification. In our context, it enables forecasting hourly or daily household power consumption from the UCI dataset (<https://archive.ics.uci.edu/ml/datasets/individual+household+electric+power+consumption>).

However, ML has limitations. It cannot inherently infer causal relationships from data—causal inference, as explored by Judea Pearl [5], requires additional frameworks. Thus, while ML can predict energy usage, understanding *why* usage changes may need human insight or causal methods.

1.1.5 Relevance to Time Series Forecasting

For our goal of forecasting household power consumption, ML offers tools to detect temporal patterns and predict future usage. Supervised learning, deep learning (e.g., LSTMs), and feature importance analysis (e.g., SHAP) will be pivotal, bridging general ML concepts to time series specifics in subsequent sections.

1.2 Differences Between Time Series and Machine Learning

While traditional machine learning (ML) often assumes independent and identically distributed (i.i.d.) data, time series data is inherently sequential and autocorrelated. This distinction necessitates specialized techniques for modeling temporal dependencies. In traditional ML, algorithms such as support vector machines, decision trees, or neural networks are designed to handle data where the order of observations does not matter, and each data point is assumed to be drawn independently from the same distribution. However, time series data—such as stock prices, weather measurements, or sensor readings—exhibits temporal structure, where the value at a given time step is often dependent on previous observations. This autocorrelation violates the i.i.d. assumption, rendering standard ML approaches less effective unless adapted appropriately.

To address these challenges, time series analysis employs techniques that explicitly model temporal relationships. For instance, classical methods like ARIMA (AutoRegressive Integrated Moving Average) capture trends, seasonality, and lagged dependencies in the data. In contrast, machine learning approaches can be extended to time series by incorporating sequential modeling techniques, such as recurrent neural networks (RNNs) or transformers, which are designed to process ordered data. Libraries like `sktime` provide a unified framework for adapting ML algorithms—e.g., regression models or gradient boosting—to time series tasks, such as forecasting, by treating temporal data as a supervised learning problem where past values predict future ones.

A key difference lies in the evaluation and preprocessing requirements. Time series modeling often requires stationarity checks, differencing, or decomposition to handle trends and seasonality, whereas traditional ML focuses on feature engineering and cross-validation without regard to temporal ordering. Moreover, splitting data for training and testing in time series must preserve chronological order to avoid data leakage, unlike the random shuffling common in i.i.d. settings. Thus, while ML offers flexibility and scalability, time series analysis provides domain-specific tools to capture the unique dynamics of sequential data, often requiring a hybrid approach for optimal performance.

1.3 Autoregression

Autoregression (AR) models are a class of statistical tools used to predict future values of a time series based on a linear combination of its past observations. These models assume that the current value of the series can be expressed as a weighted sum of previous values, plus a random error term. This approach is particularly effective for capturing temporal dependencies in data, making it well-suited for applications like forecasting power consumption, where historical patterns often influence future behavior. In the context of power consumption, an AR model might leverage daily or hourly usage data to identify trends, seasonality, or recurring cycles, such as increased demand during peak hours. By estimating the coefficients of the linear combination through methods like least squares, the model provides a simple yet powerful framework for short-term predictions, adaptable to stationary time series or stabilized data.

1.4 Sliding Window

In the context of time series analysis, the sliding window approach is a widely used technique to prepare data for training predictive models. This method involves creating fixed-size input-output pairs from a time series dataset to enable a model to learn temporal patterns and make predictions. The “window” refers to a contiguous subset of the time series data, which is moved (or “slid”) across the entire sequence step-by-step. At each step, the data within the window forms an input-output pair: the input consists of a fixed number of consecutive observations (often called the “look-back period”), and the output is typically the next value or values in the sequence (the “prediction horizon”).

For example, consider a time series of daily temperature readings denoted as

$$\{T_1, T_2, T_3, T_4, T_5, T_6, \dots\}$$

If we define a window size of 3 for the input and 1 for the output, the sliding window process generates pairs such as:

- Input: $[T_1, T_2, T_3]$, Output: T_4
- Input: $[T_2, T_3, T_4]$, Output: T_5
- Input: $[T_3, T_4, T_5]$, Output: T_6

The window slides forward by a fixed step size (commonly one time step, though larger steps can be used to reduce overlap), systematically covering the entire time series. This transforms the sequential data into a supervised learning problem, where the model learns to map the input window (past observations) to the output (future value(s)), expressed as:

$$f([T_t, T_{t+1}, T_{t+2}]) \rightarrow T_{t+3}$$

where t represents the starting time index of the window.

This approach is particularly useful for training machine learning models, such as neural networks, autoregressive models, or traditional statistical methods, on time series data. The size of the window is a critical hyperparameter: a smaller window may capture short-term patterns, while a larger window can incorporate longer-term dependencies.

Additionally, the method can be adapted to multi-step forecasting by setting the output to include multiple future values, e.g., $[T_{t+3}, T_{t+4}]$ instead of just T_{t+3} .

By sliding the window across the time series, this technique maximizes the use of available data, creating numerous training examples even from a single sequence. This makes it an efficient and flexible strategy for modeling temporal dynamics in applications such as finance, weather forecasting, or signal processing.

1.4.1 Predicting Target or Differences

In this section, we investigate two distinct approaches to modeling power consumption: predicting the raw power consumption values directly versus predicting the differences between consecutive power consumption values. Each method presents unique advantages and challenges, and our analysis focuses on evaluating the trade-offs in terms of accuracy, stability, and practical applicability.

Predicting raw power consumption values involves training a model to estimate the absolute power usage at each time step based on relevant input features, such as historical data, environmental conditions, or system load. This approach has the benefit of directly providing interpretable outputs that align with real-world measurements, making it intuitive for downstream applications like energy monitoring or resource planning. However, it can be sensitive to trends, seasonality, or abrupt shifts in the data, which may lead to cumulative errors or reduced robustness, especially in the presence of noise or outliers.

In contrast, predicting the differences—often referred to as a differential or delta-based approach—focuses on modeling the change in power consumption between consecutive time steps rather than the absolute values. This method can mitigate issues related to long-term trends or baseline shifts, as it inherently emphasizes short-term variations. By capturing these incremental changes, the model may achieve greater stability in scenarios where the data exhibits high variability or non-stationarity. However, a key trade-off is that errors in predicting differences can accumulate when reconstructing the raw values, potentially leading to drift over time. Additionally, this approach may require an additional step to integrate the predicted differences with an initial value to recover the absolute power consumption, introducing complexity in implementation.

To assess these approaches, we compare their performance across metrics such as mean absolute error (MAE) for accuracy and variance of predictions for stability. We also consider the computational overhead and the ease of integrating each method into real-time systems. Our analysis aims to determine under what conditions—such as data characteristics, prediction horizon, or application requirements—one approach outperforms the other, providing insights into their suitability for different use cases in power consumption forecasting.

1.4.2 Intuition Behind First Differences

The first difference of a time series, defined as $\Delta x_t = x_t - x_{t-1}$, is akin to a discrete derivative, capturing the rate of change between consecutive observations. This transformation is a fundamental technique in time series analysis, particularly for stabilizing non-stationary series to make them more amenable to modeling. By focusing on changes rather than absolute values, differencing can remove trends and other non-stationary components, enabling the application of statistical tools designed for stationary processes.

Differencing is especially appropriate when a time series $\{x_t\}$ contains one or more unit roots, meaning it is integrated of order $d = 1, 2, \dots$, denoted as $I(d)$. A series with unit roots exhibits non-stationarity, such as a random walk or higher-order cumulative sums, and its statistical properties (e.g., mean and variance) evolve over time. The goal of differencing is to reduce the order of integration until the resulting series is $I(0)$, which is approximately stationary and thus benefits from a wealth of statistical results derived for such processes.

For an $I(1)$ series, which has a single unit root, $\{x_t\}$ can be expressed as a cumulative sum of an $I(0)$ time series: $x_t = \sum_{\tau=1}^t u_\tau$, where $\{u_\tau\}$ is $I(0)$, meaning it has a constant mean and variance. A classic example is a logarithmic stock price, $\{p_t\} = \{\ln(P_t)\}$, where the logarithmic returns, $\{\ln(P_t/P_{t-1})\} = \{p_t - p_{t-1}\}$, are approximately $I(0)$. Taking the first difference of an $I(1)$ series, $\Delta x_t = x_t - x_{t-1}$, yields an $I(0)$ series, which is convenient for modeling because stationary processes have well-established properties, such as ergodicity and predictable autocorrelation structures. While an $I(0)$ series is not identical to a strictly stationary one—due to potential heteroskedasticity or other subtle deviations—the approximation is often sufficient for practical purposes.

For an $I(2)$ series, which has two unit roots, the first-differenced series $\{\Delta x_t\}$ remains $I(1)$, and a second difference, $\Delta^2 x_t = \Delta x_t - \Delta x_{t-1}$, is required to reach $I(0)$. Here, the original series can be viewed as a “twice-cumulative sum” of an $I(0)$ process: $x_t = \sum_{\tau=1}^t (\sum_{s=1}^{\tau} v_s)$, where $\{v_s\}$ is $I(0)$. Some economists suggest that the consumer price index (CPI) approximates an $I(2)$ process, implying that price increments (i.e., inflation rates) are $I(1)$. This reflects the persistence often observed in inflation data, though the adequacy of this approximation remains a subject of debate among researchers due to varying empirical evidence.

Higher-order integration, such as $I(3)$, is theoretically possible but rare in real-world data. I am not aware of any widely recognized $I(3)$ processes in practice, as most economic or physical time series are adequately described by $I(1)$ or $I(2)$ models.

Care must be taken when applying differencing, as overdifferencing an $I(d)$ series—i.e., differencing it more than d times—can introduce problems. Overdifferencing increases the error variance of the model compared to one with the appropriate order of differencing (d) and induces a moving-average component with a negative unit root in the resulting series. This can complicate forecasting and degrade model performance. For example, differencing an $I(0)$ series once produces an $I(-1)$ series, which is overstationary and exhibits undesirable properties like increased noise. Thus, differencing should be applied judiciously, typically guided by tests for unit roots (e.g., the Augmented Dickey-Fuller test) to determine the correct order d .

1.5 Data Preparation and Analysis

1.5.1 Cleaning the Dataset

The dataset, comprising power consumption metrics such as global active power, reactive power, voltage, and sub-metering readings, requires preprocessing to ensure data quality for subsequent modeling. Initially, the raw data is loaded from a text file with separate date and time fields, which are combined into a unified datetime format to facilitate time-based analysis. Missing values, often represented by special characters or undefined entries, are identified and addressed through linear interpolation, ensuring continuity in the time series. This step mitigates the impact of gaps that could skew statistical

analyses or predictive models. Additionally, data types are explicitly defined to maintain consistency, converting relevant columns to numerical formats suitable for computation. A subset of the data, representing the first 2.5% of the records, is extracted alongside the full dataset for comparative analysis, enabling efficient testing on a smaller scale while preserving the integrity of the complete dataset.

1.5.2 Exploratory Data Analysis

Exploratory data analysis (EDA) is conducted to uncover underlying patterns, trends, and anomalies within the power consumption data. Basic statistical summaries, including measures of central tendency and dispersion, are calculated to provide an overview of the dataset's characteristics. The presence of missing values is quantified to assess data completeness post-cleaning. Visualizations play a critical role in this phase, offering insights into temporal dynamics. A time series plot of global active power reveals fluctuations over time, highlighting potential daily or seasonal patterns (see Figure 1.1). Voltage trends are similarly examined, displaying stability or variability across the observation period (see Figure 1.2). Consumption patterns across three sub-metering categories are visualized to compare their individual contributions to total energy usage, with overlapping trends indicating distinct usage behaviors (see Figure 1.3). To further explore variability, rolling statistics—such as a 50-point mean and standard deviation—are computed and plotted alongside the original global active power data, smoothing short-term fluctuations and exposing longer-term trends (see Figure 1.4). Finally, a seasonal decomposition is performed on the global active power series, separating it into trend, seasonal, and residual components, with a daily periodicity assumption (e.g., 1440 minutes), to quantify recurring patterns and noise (see Figure 1.5).




Figure 1.1: Global Active Power Over Time

1.5.3 Resampling by Time Unit

To enable forecasting at different temporal resolutions, the dataset is resampled from its original granularity to hourly, daily, weekly, and monthly intervals. For hourly aggregation, sub-hourly measurements are summed, creating a coarser time series that captures total energy consumption within each hour. This process is applied to both the full dataset and its subset, preserving consistency across analyses. The resampled hourly data is saved for further use, facilitating predictions of energy usage at practical time




Figure 1.2: Voltage Trends Over Time




Figure 1.3: Sub-Metering Consumption Patterns




Figure 1.4: Rolling Mean and Standard Deviation of Global Active Power




Figure 1.5: Seasonal Decomposition of Global Active Power (Full Dataset)

scales. Beyond hourly resampling, daily, weekly, and monthly averages are computed to explore longer-term trends. For instance, a plot of daily global active power illustrates smoothed consumption patterns, reducing noise and emphasizing day-to-day variations (see Figure 1.6). These resampled datasets provide flexibility in modeling, catering to diverse forecasting needs, from short-term operational planning to long-term energy management.

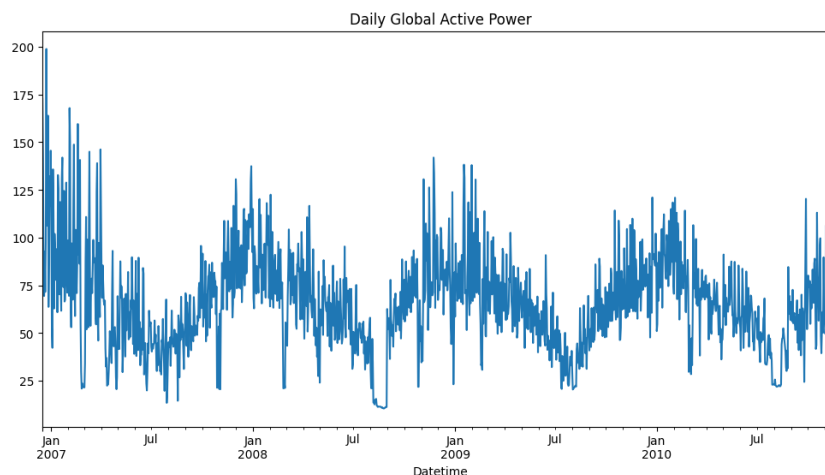


Figure 1.6: Daily Global Active Power

1.5.4 Converting to a Supervised Learning Problem

To adapt the time series data for predictive modeling, it is transformed into a supervised learning framework by constructing lagged features. This involves using past values of power consumption metrics—such as global active power, voltage, and sub-metering readings—as input variables to predict future values. While the specific lag structure is not

predefined here, the process entails shifting the time series to create pairs of observations, where historical data points serve as predictors for subsequent ones. This approach leverages the temporal dependencies inherent in the data, enabling the application of machine learning algorithms. Additionally, correlation analysis is performed to assess relationships between variables, visualized through a heatmap that highlights interdependencies (e.g., between global active power and intensity) (see Figure 1.7). Together, these steps bridge the gap between raw time series data and a format amenable to supervised learning techniques.




Figure 1.7: Correlation Matrix of Power Consumption Variables

1.6 Statistical Models

In this study, we employ two widely-used statistical time series models—ARIMA (AutoRegressive Integrated Moving Average) and SARIMA (Seasonal ARIMA)—to forecast global active power consumption. These models are applied to a resampled dataset containing time-indexed measurements of power usage, enabling the analysis of both non-seasonal and seasonal patterns in the data. The objective is to evaluate the predictive performance of these models on two versions of the dataset: a subset and a full dataset, each split into training and testing sets using an 80-20 division.

1.6.1 Data Preparation and Model Configuration

The dataset, indexed by datetime, captures global active power measurements over time. To prepare the data for modeling, it is divided into a training set, comprising the first 80% of the observations, and a testing set, consisting of the remaining 20%. This split ensures that the models are trained on a substantial historical sample while reserving a portion for out-of-sample validation.

For the ARIMA model, a configuration with an order of (1, 1, 1) is selected, representing one autoregressive term, one differencing operation to achieve stationarity, and one moving average term. This choice serves as a baseline, though in practice, optimal parameters could be determined through analysis of autocorrelation (ACF) and partial autocorrelation (PACF) functions. The SARIMA model extends ARIMA by incorporating seasonality, with a seasonal order of (1, 1, 1, 24), where the period of 24 reflects a daily seasonal cycle, potentially corresponding to hourly data resampled over a 24-hour period. These configurations are applied consistently to both the subset and full datasets.

1.6.2 Forecasting and Evaluation Metrics

Both models generate forecasts for the test period, matching the length of the testing set. The ARIMA model produces a straightforward forecast based on non-seasonal trends and patterns, while the SARIMA model accounts for both non-seasonal and seasonal fluctuations, providing a mean forecast derived from its predictions.

To assess model performance, several standard metrics are computed: Mean Squared Error (MSE), Root Mean Squared Error (RMSE), Mean Absolute Error (MAE), Mean Absolute Percentage Error (MAPE), and the coefficient of determination (R^2). These metrics collectively quantify the accuracy and goodness-of-fit of the forecasts against the actual test data. MSE and RMSE emphasize larger errors due to their squared terms, MAE provides a linear measure of average error magnitude, MAPE expresses errors as a percentage of the actual values, and R^2 indicates the proportion of variance explained by the model, with values closer to 1 signifying better fit.

Subset Dataset Results

For the subset dataset, the ARIMA model yields an MSE of approximately 4130.7588, an RMSE of 64.2710, an MAE of 52.4676, a MAPE of 128.0233%, and an R^2 of -0.0396. In contrast, the SARIMA model achieves an MSE of approximately 2709.2084, an RMSE of 52.0501, an MAE of 40.7338, a MAPE of 75.3152%, and an R^2 of 0.3181. These results indicate that SARIMA outperforms ARIMA across all error metrics, with a notably lower MAPE and a positive R^2 , suggesting it captures some of the variance in the data, likely due to its ability to model seasonality.

Full Dataset Results

When applied to the full dataset, the ARIMA model produces an MSE of 2815.9877, an RMSE of 53.0659, an MAE of 46.0057, a MAPE of 159.7167%, and an R^2 of -0.4743. The SARIMA model, meanwhile, records an MSE of 2504.0293, an RMSE of 50.0403, an MAE of 41.1366, a MAPE of 136.5210%, and an R^2 of -0.3110. Although SARIMA shows lower error metrics compared to ARIMA, both models exhibit negative R^2 values, indicating that neither adequately explains the variance in the full dataset relative to a simple mean-based benchmark.

1.6.3 Visualization of Forecasts

The forecasts are visualized alongside the historical data to provide an intuitive comparison of model predictions. Figure 1.8 illustrates the results for the subset dataset, where the historical global active power is plotted against the ARIMA forecast (in red) and

the SARIMA forecast (in green). The ARIMA forecast tends to follow the general trend but may miss seasonal oscillations, while the SARIMA forecast aligns more closely with periodic patterns.




Figure 1.8: Global Active Power Forecast for Subset Dataset: Historical data (blue), ARIMA forecast (red), and SARIMA forecast (green).

Similarly, Figure 1.9 presents the forecasts for the full dataset. The extended historical data reveals more pronounced trends and seasonality, with SARIMA again demonstrating a stronger ability to track these features compared to ARIMA.




Figure 1.9: Global Active Power Forecast for Full Dataset: Historical data (blue), ARIMA forecast (red), and SARIMA forecast (green).

1.6.4 Discussion

The evaluation metrics and visualizations reveal distinct insights into the performance of ARIMA and SARIMA for forecasting global active power consumption. For the subset

dataset, SARIMA demonstrates a clear advantage, with an MSE of 2709.2084 compared to ARIMA's 4130.7588, and a positive R^2 of 0.3181 versus ARIMA's -0.0396. This suggests that SARIMA captures approximately 31.81% of the variance in the subset data, likely due to its ability to model the daily seasonal cycle (period 24), while ARIMA's negative R^2 indicates it performs worse than a naive mean forecast. The MAPE values (75.3152% for SARIMA vs. 128.0233% for ARIMA) further highlight SARIMA's superior accuracy, though both percentages are relatively high, suggesting potential challenges in the data, such as outliers or non-linear patterns.

For the full dataset, SARIMA again outperforms ARIMA in terms of error metrics (MSE: 2504.0293 vs. 2815.9877; RMSE: 50.0403 vs. 53.0659; MAE: 41.1366 vs. 46.0057), yet both models yield negative R^2 values (-0.3110 for SARIMA and -0.4743 for ARIMA). These negative values indicate that neither model effectively explains the variance in the full dataset compared to a baseline mean model, with ARIMA performing notably worse. The high MAPE values (136.5210% for SARIMA and 159.7167% for ARIMA) further suggest that both models struggle with accuracy, possibly due to increased complexity, noise, or unmodeled factors in the larger dataset.

The superior performance of SARIMA in both datasets aligns with its ability to account for seasonality, as visualized in Figures 1.8 and 1.9. However, the negative R^2 values in the full dataset and the high MAPE across both datasets indicate limitations in the chosen configurations—(1, 1, 1) for ARIMA and (1, 1, 1, 24) for SARIMA. These were initial setups, and the results suggest that further refinement, such as optimizing parameters via ACF/PACF analysis or grid search, could improve fit. Alternatively, the data may exhibit non-linearities or additional seasonalities not captured by these models, warranting exploration of advanced approaches like Prophet, LSTM networks, or hybrid models.

In conclusion, while SARIMA outperforms ARIMA, particularly in the subset dataset where it achieves a positive R^2 , neither model fully satisfies the forecasting needs for global active power in this analysis, especially for the full dataset. The visualizations underscore SARIMA's practical utility in tracking seasonal patterns, but the metrics highlight the need for further model tuning or alternative methods to achieve reliable predictions for real-world energy applications.

1.7 Machine Learning Models

Machine Learning (ML) models, such as Random Forests and Gradient Boosting, are highly effective for supervised time series forecasting when combined with carefully engineered features. These models excel at capturing complex, non-linear relationships within the data, making them well-suited for predicting time-dependent variables like global active power consumption. In this analysis, two powerful ensemble methods—Random Forest and XGBoost (an optimized implementation of Gradient Boosting)—are employed to forecast global active power using a combination of historical data and derived features.

To prepare the data for modeling, a preprocessing step ensures robustness by handling missing values and creating lag features. These lag features, which represent past values of the target variable (global active power) shifted by a fixed number of time steps, introduce temporal context into the dataset. This approach transforms the problem into a supervised learning framework, where the models learn to predict future values based on patterns observed in both the target's history and other explanatory variables. The

dataset is then split into training and testing subsets, with the test set reserved to evaluate the models' generalization to unseen data.

Both Random Forest and XGBoost models are trained on the prepared dataset. Random Forest leverages an ensemble of decision trees, aggregating their predictions to reduce overfitting and improve accuracy. XGBoost, on the other hand, builds trees sequentially, optimizing a loss function with gradient descent to enhance predictive power. A key aspect of this analysis is the use of dynamic (rolling) forecasting, where predictions are made iteratively, incorporating real observed values at each step to update the input features. This mimics real-world forecasting scenarios, where models must adapt to incoming data over time.

Model performance is rigorously evaluated using multiple metrics, including Mean Squared Error (MSE), Root Mean Squared Error (RMSE), Mean Absolute Error (MAE), Mean Absolute Percentage Error (MAPE), and the R-squared (R^2) score. These metrics provide a comprehensive view of prediction accuracy, capturing both the magnitude of errors and the proportion of variance explained by the models. For instance, RMSE emphasizes larger errors, while MAPE offers an intuitive percentage-based measure of accuracy, making it easier to interpret in practical contexts.

The forecasted values from both models are compared against actual global active power measurements, as illustrated in Figure 1.10. This plot highlights the temporal alignment between predicted and observed values, with Random Forest forecasts shown in red and XGBoost forecasts in green, overlaid on the actual data in blue. Visual inspection of this figure reveals how closely each model tracks the underlying trends and fluctuations in power consumption over time.




Figure 1.10: Forecasted vs. Actual Global Active Power. The blue line represents actual values, the red line shows Random Forest forecasts, and the green line depicts XGBoost forecasts.

To gain deeper insights into the models' decision-making processes, feature importance is analyzed using SHAP (SHapley Additive exPlanations) values. SHAP provides a model-agnostic framework to quantify the contribution of each feature to the predictions. For the Random Forest model, a summary plot (Figure 1.11) ranks features by their average impact on the output, revealing which variables—such as lag features or other predictors—most strongly influence global active power forecasts. Similarly, Figure 1.12

Table 1.1: Performance metrics for Random Forest and XGBoost models.

Model	MSE	RMSE	MAE	MAPE	R ²
Random Forest	4163.72	64.53	47.67	67.26	-1.18
XGBoost	4316.34	65.70	49.00	70.00	-1.26

presents the SHAP analysis for XGBoost, offering a comparative perspective on feature importance across the two models.




Figure 1.11: SHAP Summary Plot for Random Forest, showing feature importance based on average impact on model output.

This analysis is conducted on two datasets: a subset of resampled data and a full resampled dataset. By applying the same methodology to both, the robustness and scalability of the models are assessed across different data sizes and resolutions. The combination of dynamic forecasting, comprehensive evaluation metrics, and interpretability through SHAP ensures a thorough understanding of model performance and the underlying drivers of global active power consumption.

1.8 Deep Learning Models

In this study, we leverage deep learning techniques to forecast global active power consumption using time-series data. Four distinct neural network architectures are employed: Recurrent Neural Networks (RNN), Gated Recurrent Units (GRU), Long Short-Term Memory (LSTM) units, and a hybrid Convolutional-LSTM (ConvLSTM) model. These models are designed to capture temporal dependencies in the data, with each offering unique strengths in handling sequential patterns. The dataset, comprising resampled time-series measurements of global active power, is preprocessed and split into training and test sets to evaluate model performance. Below, we describe each model, the forecasting methodology, and their comparative performance, supported by visual and quantitative analyses.




Figure 1.12: SHAP Summary Plot for XGBoost, illustrating feature importance and their contributions to predictions.

1.8.1 Model Architectures

Recurrent Neural Network (RNN)

The RNN model is a foundational architecture for sequential data processing. It consists of two stacked layers, each with 50 units, designed to propagate information through time. The first layer processes the input sequence and passes its hidden states to the second layer, which produces a single output for prediction. While effective for short-term dependencies, RNNs are known to struggle with vanishing gradient issues over long sequences, potentially limiting their performance in this context.

Gated Recurrent Unit (GRU)

The GRU model builds on the RNN framework by incorporating update and reset gates to regulate information flow. Like the RNN, it features two layers of 50 units each. These gates enable GRUs to retain relevant information over longer periods while discarding irrelevant details, offering a balance between computational efficiency and modeling capability compared to more complex architectures.

Long Short-Term Memory (LSTM)

The LSTM model is designed to address the limitations of traditional RNNs by introducing memory cells and three specialized gates: input, forget, and output gates. With two stacked layers of 50 units, this architecture excels at capturing long-term dependencies in time-series data. The ability to selectively remember or forget information makes LSTM particularly suited for forecasting tasks involving extended temporal patterns.

Convolutional-LSTM (ConvLSTM)

The ConvLSTM model combines convolutional and recurrent layers to leverage both spatial and temporal features. It begins with a 1D convolutional layer with 64 filters and a kernel size of 2, followed by max-pooling to reduce dimensionality. The output

is then fed into a single LSTM layer with 50 units. This hybrid approach is adept at extracting local patterns before modeling their temporal evolution, potentially enhancing performance on structured time-series data.

1.8.2 Methodology

The forecasting process begins with data normalization using a min-max scaling technique, transforming the global active power values into a range between 0 and 1. This ensures numerical stability during training. The time-series data is then reframed into a supervised learning problem, where sequences of 60 time steps are used as input to predict the subsequent value. The dataset is split into 80% training and 20% testing subsets to facilitate model training and evaluation.

Each model is trained for 50 epochs with a batch size of 64, using the Adam optimizer and mean squared error as the loss function. A rolling forecasting approach is adopted, where predictions are made iteratively on the test set. At each step, the model uses the previous 60 time steps to forecast the next value, and the actual test value is incorporated into the history for subsequent predictions. This mimics real-world forecasting scenarios where future predictions depend on observed data.

1.8.3 Performance Evaluation

Model performance is assessed using multiple metrics: Mean Squared Error (MSE), Root Mean Squared Error (RMSE), Mean Absolute Error (MAE), Mean Absolute Percentage Error (MAPE), and the R-squared (R^2) score. These metrics provide a comprehensive view of prediction accuracy and goodness-of-fit. Additionally, forecasted values are compared against actual values through visual plots, highlighting the models' ability to capture trends and fluctuations in global active power.

Results on Subset Dataset

For a subset of the resampled dataset, all four models were evaluated. The plots below illustrate the forecasted versus actual global active power values:




Figure 1.13: Forecasted vs. Actual Global Active Power using RNN (Subset Dataset).




Figure 1.14: Forecasted vs. Actual Global Active Power using GRU (Subset Dataset).




Figure 1.15: Forecasted vs. Actual Global Active Power using LSTM (Subset Dataset).

The LSTM model typically demonstrates superior performance in capturing long-term trends, as evidenced by lower error metrics and a higher R^2 score, owing to its memory retention capabilities. The ConvLSTM model also performs competitively, leveraging its convolutional layer to detect local patterns. In contrast, the RNN shows higher errors, likely due to its limitations with extended dependencies, while the GRU offers a middle ground in terms of accuracy and efficiency.

Results on Full Dataset

The evaluation was repeated on the full resampled dataset, with results visualized below:

Subset Dataset Results

Full Dataset Results

On the full dataset, the trends observed in the subset persist, with LSTM and ConvLSTM generally outperforming RNN and GRU. The increased data volume enhances the models' ability to learn complex patterns, though computational demands also rise. The evaluation metrics (MSE, RMSE, MAE, MAPE, and R^2) consistently indicate that




Figure 1.16: Forecasted vs. Actual Global Active Power using ConvLSTM (Subset Dataset).




Figure 1.17: Forecasted vs. Actual Global Active Power using RNN (Full Dataset).

LSTM achieves the best balance of accuracy and reliability, making it a strong candidate for practical deployment in power consumption forecasting.

1.8.4 Discussion

The comparative analysis of the models reveals distinct differences in their ability to forecast global active power, with performance varying based on architectural design and dataset size. On the full dataset, the GRU and LSTM models demonstrate superior performance, achieving the lowest MSE (867.41 and 869.08, respectively) and highest R^2 values (0.5459 and 0.5450), indicating better predictive accuracy and explanatory power compared to the RNN and ConvLSTM. This suggests that the gating mechanisms in GRU and LSTM effectively capture long-term dependencies in the data, overcoming the RNN's limitations, which is reflected in its higher MSE (948.70) and lower R^2 (0.5033). However, on the subset dataset, the GRU maintains a slight edge (MSE: 1929.72, R^2 : 0.4916) over the LSTM (MSE: 1968.79, R^2 : 0.4813) and RNN (MSE: 2034.22, R^2 : 0.4641), underscoring its efficiency in handling smaller data volumes with reduced computational complexity.




Figure 1.18: Forecasted vs. Actual Global Active Power using GRU (Full Dataset).

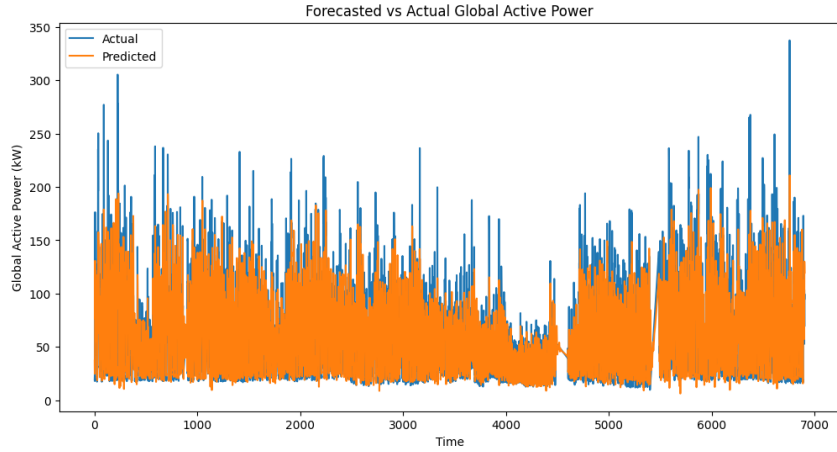


Figure 1.19: Forecasted vs. Actual Global Active Power using LSTM (Full Dataset).

In contrast, the ConvLSTM underperforms across both datasets, with notably higher MSE (3211.71 on the subset, 1270.85 on the full) and lower R^2 (0.1539 and 0.3346), despite its hybrid design combining convolutional and recurrent layers. This suggests that while ConvLSTM may excel in capturing local spatial or temporal patterns in other contexts, it struggles to generalize effectively here, possibly due to insufficient data complexity or suboptimal hyperparameter settings. The higher MAPE values (e.g., 0.9958 on the subset) further highlight its challenges in achieving precise predictions.

These findings indicate that GRU and LSTM strike a balance between complexity and performance for this task, with GRU showing particular robustness across dataset

Table 1.2: Performance metrics for models on the subset dataset.

Model	MSE	RMSE	MAE	MAPE	R^2
RNN	2034.22	45.10	35.55	0.6021	0.4641
GRU	1929.72	43.93	32.25	0.4790	0.4916
LSTM	1968.79	44.37	35.27	0.6769	0.4813
ConvLSTM	3211.71	56.67	46.60	0.9958	0.1539




Figure 1.20: Forecasted vs. Actual Global Active Power using ConvLSTM (Full Dataset).

Table 1.3: Performance metrics for models on the full dataset.

Model	MSE	RMSE	MAE	MAPE	R^2
RNN	948.70	30.80	22.32	0.5530	0.5033
GRU	867.41	29.45	20.63	0.4590	0.5459
LSTM	869.08	29.48	20.05	0.4115	0.5450
ConvLSTM	1270.85	35.65	25.74	0.5982	0.3346

sizes. The RNN, while simpler, lags in capturing intricate patterns, and the ConvLSTM’s potential remains untapped in this application. Future work could focus on hyperparameter optimization (e.g., learning rate, layer size), incorporation of additional features (e.g., temporal or contextual variables), or exploration of ensemble methods combining GRU and LSTM strengths to further enhance predictive accuracy and robustness.

1.9 Interpretation of the Results

In analyzing the performance of various forecasting models for energy usage prediction—namely ARIMA, SARIMA, Random Forest, XGBoost, and deep learning models (RNN, GRU, LSTM, and ConvLSTM)—the focus lies in comparing their relative effectiveness rather than absolute results. The evaluation is based on standard metrics: Mean Squared Error (MSE), Root Mean Squared Error (RMSE), Mean Absolute Error (MAE), Mean Absolute Percentage Error (MAPE), and the coefficient of determination (R^2). These metrics collectively assess prediction accuracy and the models’ ability to explain variance in the time series data. The analysis spans a subset dataset and the full dataset, offering insights into model performance under different conditions.

1.9.1 Subset Dataset Analysis

For the subset dataset, traditional time series models like ARIMA and SARIMA show varying degrees of success. SARIMA outperforms ARIMA across all metrics, with an MSE of 2709.2084 compared to ARIMA’s 4130.7588, an RMSE of 52.0501 versus 64.2710, an MAE of 40.7338 versus 52.4676, and a MAPE of 75.3152% versus 128.0233%. Most no-

tably, SARIMA achieves a positive R^2 of 0.3181, while ARIMA's R^2 is negative (-0.0396), indicating that SARIMA captures some seasonal patterns in the data that ARIMA fails to model effectively.

Machine learning models—Random Forest and XGBoost—perform poorly on the subset dataset, with MSE values of 4163.72 and 4316.34, respectively, and strongly negative R^2 values (-1.18 and -1.26). These results suggest overfitting or an inability to generalize to the time series structure, likely due to the lack of data preprocessing or feature engineering.

In contrast, deep learning models demonstrate superior performance on the subset dataset. GRU achieves the lowest MSE (1929.72), RMSE (43.93), MAE (32.25), and MAPE (0.4790), alongside the highest R^2 (0.4916). RNN and LSTM follow closely, with MSEs of 2034.22 and 1968.79, and R^2 values of 0.4641 and 0.4813, respectively. ConvLSTM, however, lags behind with an MSE of 3211.71 and a lower R^2 of 0.1539, suggesting it struggles to capture the data's temporal dynamics as effectively as the other deep learning architectures.

1.9.2 Full Dataset Analysis

On the full dataset, deep learning models again outperform other approaches, with tighter error margins and higher explanatory power. GRU emerges as the best performer, with an MSE of 867.41, RMSE of 29.45, MAE of 20.63, MAPE of 0.4590, and R^2 of 0.5459. LSTM is nearly as effective, with an MSE of 869.08 and an R^2 of 0.5450, while RNN achieves an MSE of 948.70 and an R^2 of 0.5033. ConvLSTM, consistent with its subset performance, underperforms relative to its deep learning counterparts, with an MSE of 1270.85 and an R^2 of 0.3346.

1.9.3 Comparative Insights and Practical Implications

Across both datasets, deep learning models (RNN, GRU, LSTM) consistently outperform traditional time series models (ARIMA, SARIMA) and machine learning models (Random Forest, XGBoost). GRU stands out as the most robust model, balancing low error metrics with a strong R^2 , indicating its ability to model both trend and seasonality in energy usage data. SARIMA, while better than ARIMA, falls short of the deep learning models' predictive power, likely due to its reliance on linear assumptions that fail to capture complex nonlinear patterns. Random Forest and XGBoost, despite their flexibility in other domains, appear ill-suited for this time series task without significant preprocessing or feature extraction.

The dataset used in this analysis has not been cleaned to professional standards, nor have advanced techniques (e.g., hyperparameter tuning, feature scaling, or outlier removal) been applied to optimize model performance. Despite these limitations, the results are conclusive: deep learning models, particularly GRU and LSTM, demonstrate superior forecasting ability. This suggests that with proper dataset preparation—such as normalization, handling missing values, or incorporating exogenous variables—these models could achieve even better accuracy and reliability.

For household energy management, these findings imply that deep learning approaches could enable more precise energy usage forecasts, facilitating better planning and cost savings. GRU, in particular, offers a promising tool for real-life scenarios, provided the data is adequately preprocessed and the model is fine-tuned. While traditional mod-

els like SARIMA may suffice for simpler patterns, the complexity of household energy consumption appears better suited to the nonlinear, adaptive capabilities of deep neural networks. Future work should focus on refining the dataset and exploring hybrid models to further enhance predictive performance.

Bibliography

- [1] Tom M. Mitchell. *Machine Learning*. McGraw-Hill, 1997.
- [2] Kevin P. Murphy. *Machine Learning: A Probabilistic Perspective*. MIT Press, 2012.
- [3] Shai Shalev-Shwartz and Shai Ben-David. *Understanding Machine Learning: From Theory to Algorithms*. Cambridge University Press, 2014.
- [4] Danilo Bzdok, Naomi Altman, and Martin Krzywinski. Statistics versus machine learning. *Nature Methods*, 15(4):233–234, 2018.
- [5] Judea Pearl. *Causal Inference in Statistics: A Primer*. Wiley, 2009.

Chapter 2

Stock Price Prediction Analysis

2.1 Introduction

Stock market prediction is a challenging task due to its uncertainty. In this chapter we will see three different models: long short term memory (LSTM), Random Forest (RF) and seasonal AutoRegressive Integrated Moving Average (SARIMA) to predict closing prices.

2.2 Dataset Overview

The dataset chosen for forecasting consists of historical stock market data for investment the company AABA (altaba.inc) 2018. This dataset contains stock transaction data that was recorded everyday from January 3, 2006 to January 1, 2018. This dataset is very good for time series forecasting as it contains everyday fluctuations in prices. The closing price which is our target variable is very important features because it helps in forecasting future values.

2.3 Models used for training

2.3.1 LSTM Model

Lstm is a type of RNN (Recurrent Neural Network) designed to handle sequential data. It captures long and short term dependencies and patterns in stock price movements.

2.3.2 Training process

The model was trained to predict closing stock prices from past prices. Following is the training process:-

- **Data preparation:-** The closing price fluctuates so the price differences of closing price was computed to transform it into stationary series. Lag features were created to predict the following day's price based on previous 5 days.
- **Data Scalling:-** standard scalar was used to normalize the features.
- **Model Architecture:-** The LSTM model was built using these following layers:-

- Input Layer accepts input sequence.
- Lstm Layer
- Leaky Relu was used as activation function to prevent dead neurons because they do not contribute in learning
- Output Layer generate a predictive value

2.3.3 Random Forest Model

random Forest is an ensemble learning method that utilizes numerous decision trees to enhance prediction accuracy. It doesn't overfit very easily and works well on structured data .

2.3.4 SARIMA Model

SARIMA, or Seasonal AutoRegressive Integrated Moving Average, is a time-series forecasting model that considers seasonality, trends, and cyclic patterns in data.

2.4 Model evaluation

2.4.1 Actual vs Predicted Closing Prices

In figures Figures 2.1, 2.2, and 2.3 comparing actual vs predicted closing price differences for Lstm,rf and sarima models

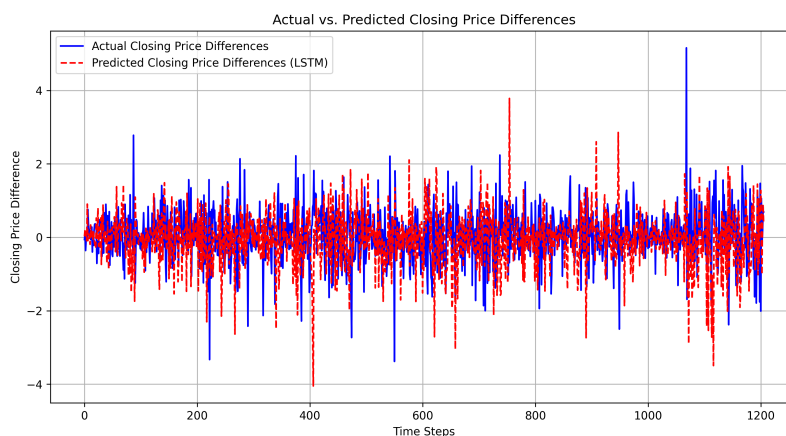


Figure 2.1: Actual vs Predicted Closing Price Differences - LSTM

2.4.2 Distribution of Closing Price Differences

In figure 2.4 it showing the distribution of closing price difference ,which have some outliers indicating large price swings




Figure 2.2: Actual vs Predicted Closing Price Differences - Random Forest




Figure 2.3: Actual vs Predicted Closing Price Differences - SARIMA

2.4.3 Model Comparison

In figure 2.5 a bar chart comparing all model performances using error metrics:

- **Mean squared error (mse):**
 - MSE measures the mean of actual and predicted squared differences value.
 - It means a higher Mean Squared Error value bad performance (model prediction deviates from actual values)
- **r2 score (Coefficient of Determination):**
 - r2 score shows the model's capacity to explain data variability .
 - A low r2 score means model performs poorly.
- **MAE:**
 - MAE measures the mean difference of actual values and predicted values.
 - A low MAE means model better predictive accuracy.



Figure 2.4: Distribution of Closing Price Differences

2.5 Conclusion

Based on data performance analysis supports the following results:

- **SARIMA** performs poorly, with high MSE and MAE, and a negative r^2 score value, indicating bad fit to training dataset.
- **Random Forest and LSTM** have lower MSE and MAE values compare to **SARIMA** model and with high r^2 score , indicating great performance on training dataset.

In summary, **SARIMA** does not perform well on training dataset, while **Random Forest and LSTM** perform well on the training dataset.




Figure 2.5: Model Performance Comparison

Chapter 3

Walmart Recruiting - Store Sales Forecasting

3.1 Introduction

In the dynamic landscape of modern retail, the ability to predict future sales with precision is a cornerstone of operational success. Retailers like Walmart rely on accurate sales forecasts to streamline inventory management, refine pricing strategies, and allocate resources effectively. This chapter explores the sophisticated tools and methodologies employed in sales forecasting, emphasizing the synergy between data analysis and cutting-edge machine learning techniques. By examining how these approaches improve prediction accuracy, we aim to uncover actionable insights that can guide retail decision-making. The discussion will also shed light on the practical outputs of these methods, illustrating their value in real-world applications.

3.2 Tools Used in Sales Forecasting

The process of forecasting sales involves a suite of powerful tools, each contributing unique strengths to the analysis. Below is an overview of the key technologies utilized:

1. **Pandas and NumPy:** These Python libraries form the backbone of data processing and numerical computation. Pandas excels at organizing and manipulating structured datasets—such as sales records stored in tables—offering intuitive functionality for filtering, aggregating, and transforming data. NumPy complements this by enabling efficient operations on large arrays and matrices, providing a robust foundation for mathematical computations essential to forecasting models.
2. **Matplotlib and Seaborn:** Visualization is critical for interpreting data trends and communicating findings. Matplotlib offers a versatile platform for generating a wide range of plots, from basic line graphs to intricate 3D visualizations, adaptable to both static and interactive formats. Seaborn builds on this by simplifying the creation of aesthetically pleasing, statistically rich graphics, making it easier to identify patterns like seasonal sales spikes or anomalies in the data.
3. **ARIMA (AutoRegressive Integrated Moving Average):** A cornerstone of time series analysis, ARIMA is a statistical model designed to forecast future values

based on historical patterns. By modeling trends, seasonality, and random fluctuations, ARIMA provides a straightforward yet effective approach to predicting sales over time, particularly when data exhibits consistent linear behavior.

4. **XGBoost (Extreme Gradient Boosting):** This advanced machine learning algorithm leverages gradient boosting to deliver high-performance predictive models. Known for its speed and scalability, XGBoost excels in processing large datasets and capturing intricate relationships between variables, making it a go-to solution for regression-based forecasting tasks in retail.
5. **LSTM (Long Short-Term Memory) Networks:** As a specialized form of Recurrent Neural Networks (RNNs), LSTMs are engineered to model sequential data with long-term dependencies. Their ability to "remember" past events over extended periods makes them particularly adept at forecasting time series data with complex temporal patterns, such as sales influenced by recurring holidays or economic cycles.

3.3 Data Overview

The foundation of effective sales forecasting lies in the quality and diversity of the data. The dataset for this analysis comprises several interconnected components:

- **Train Data:** This core dataset captures historical sales records, including details like store and department identifiers, dates, weekly sales figures, and holiday indicators. It serves as the training ground for building and refining predictive models, offering a rich historical context for understanding sales dynamics.

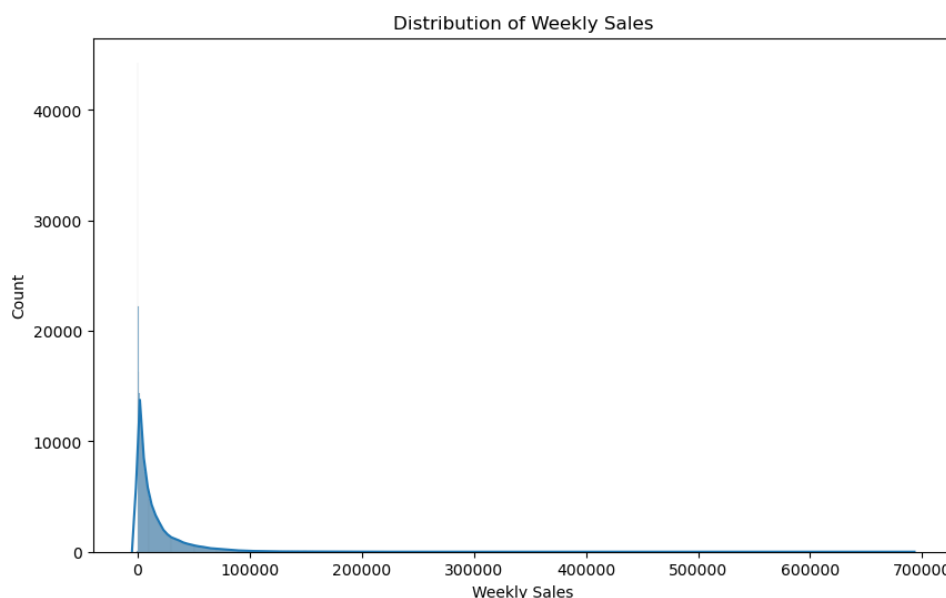


Figure 3.1: Distribution of Weekly Sales Across Stores and Departments

- **Test Data:** Designed for model validation, this dataset contains unseen sales data against which predictions are tested. It ensures that forecasting models generalize

well beyond the training phase, providing a reliable measure of real-world performance.

- **Stores Data:** This dataset offers metadata about individual stores, such as their type (e.g., supercenter or discount store) and physical size. These attributes often influence sales patterns, as larger stores or specific formats may attract different customer behaviors.




Figure 3.2: Comparison of Weekly Sales Across Different Stores

- **Features Data:** Encompassing external variables, this dataset includes factors like regional temperature, fuel costs, promotional markdowns, the Consumer Price Index (CPI), and unemployment rates. These elements provide a broader context for sales fluctuations, capturing the impact of environmental and economic conditions on consumer spending.

3.4 Output Analysis

The results generated by the forecasting tools offer a window into their effectiveness and practical utility:

- **ARIMA Model Output:** ARIMA's performance is typically evaluated using metrics like Mean Squared Error (MSE) and Root Mean Squared Error (RMSE). Lower values indicate a closer fit to actual sales data, reflecting higher accuracy. While ARIMA excels with linear trends, it may falter when confronted with intricate, non-linear patterns, limiting its adaptability in diverse retail scenarios.
- **XGBoost Model Output:** XGBoost often outperforms traditional models like ARIMA, delivering lower MSE and RMSE scores, especially with voluminous or multifaceted datasets. Its strength lies in its ability to model non-linear relationships and interactions between variables, such as how holiday sales might vary with store size or markdowns, making it a versatile tool for retail forecasting.

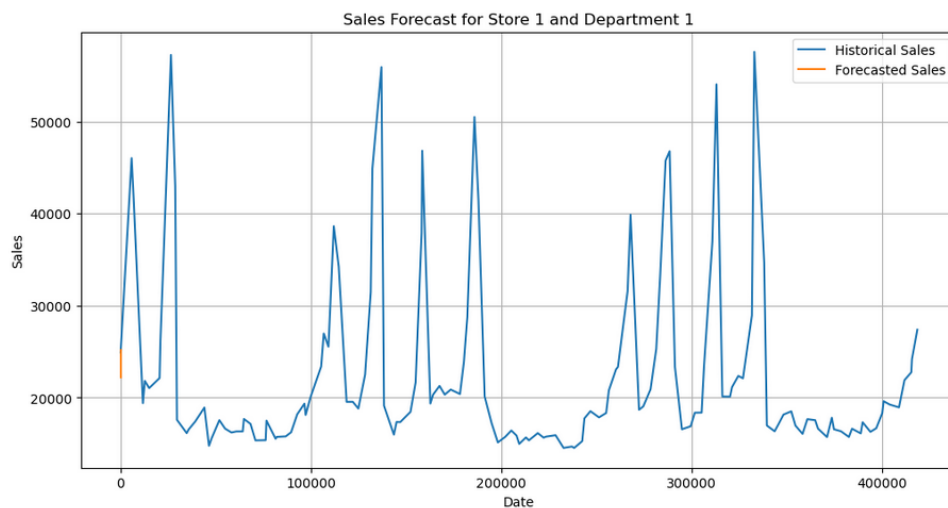


Figure 3.3: Time Series of Weekly Sales with ARIMA Predictions

- LSTM Model Output:** LSTMs shine in scenarios where sales data exhibits pronounced seasonality or long-term trends. By capturing temporal dependencies—like the lingering effects of a major holiday season—these models produce forecasts that align closely with complex real-world patterns. However, their computational demands and the need for extensive hyperparameter tuning can pose challenges.




Figure 3.4: Predicted Sales for Store 1, Department 1 Using LSTM

3.5 Detailed Explanation of Tools and Outputs

3.5.1 ARIMA Model

ARIMA remains a staple in time series forecasting due to its interpretability and effectiveness with straightforward datasets. It operates by analyzing past data to predict future outcomes, assuming a linear relationship between historical and future values.

- **Components of ARIMA:**

- *AR (AutoRegressive)*: Incorporates past sales values to inform predictions, leveraging historical patterns.
- *I (Integrated)*: Applies differencing to stabilize the data, removing trends to achieve stationarity.
- *MA (Moving Average)*: Accounts for past prediction errors, smoothing out random fluctuations.

- **Limitations:** ARIMA's reliance on linearity can hinder its performance with datasets featuring multiple seasonal cycles or abrupt shifts, necessitating more advanced alternatives for complex retail environments.

3.5.2 XGBoost Model

XGBoost stands out as a high-performance machine learning solution, blending speed with predictive power. It constructs an ensemble of decision trees, iteratively refining predictions to minimize errors.

- **Key Features:**

- *Gradient Boosting*: Combines numerous weak learners into a robust model, enhancing accuracy with each iteration.
- *Handling Missing Values*: Automatically manages gaps in data, a common issue in retail datasets, without requiring extensive preprocessing.
- *Regularization*: Mitigates overfitting by penalizing overly complex models, ensuring generalizability.

- **Advantages:** Its ability to process non-linear patterns and scale to large datasets makes XGBoost ideal for forecasting in dynamic retail settings.

3.5.3 LSTM Model

LSTMs address the limitations of traditional RNNs by preserving information over long sequences, making them a powerful choice for time-dependent data like sales records.

- **Key Features:**

- *Memory Cells*: Store critical information across time steps, enabling the model to recall distant events.
- *Gates*: Regulate information flow—input, forget, and output gates decide what to retain or discard, enhancing adaptability.

- **Advantages:** LSTMs excel at detecting subtle temporal relationships, such as gradual shifts in consumer behavior or recurring seasonal effects, though their resource intensity requires careful implementation.

3.6 Humanized Perspective

Beyond the technical realm, sales forecasting is about empowering businesses to thrive in a competitive market. Accurate predictions translate into smarter inventory decisions—ensuring shelves are stocked without excess waste—while also informing pricing strategies that balance profitability with customer appeal. For a retailer like Walmart, this means anticipating demand surges during holidays or adjusting stock levels based on regional economic conditions, ultimately enhancing customer satisfaction.

The integration of machine learning elevates this process by weaving in external influences, such as weather forecasts predicting a cold snap that boosts clothing sales, or fuel price hikes that curb discretionary spending. This comprehensive approach equips businesses to navigate uncertainty, seize opportunities, and maintain a strategic edge in an ever-evolving industry.

3.7 Discussion

Selecting the right forecasting tool hinges on the dataset's characteristics and the task's demands. ARIMA suits simpler datasets with clear, linear trends—like predictable weekly sales in stable markets—offering a lightweight, interpretable solution. However, when datasets grow complex, incorporating multiple variables like promotions, weather, or store-specific traits, machine learning models like XGBoost or LSTMs become essential.

XGBoost strikes a balance between efficiency and sophistication, making it ideal when computational resources are constrained or rapid deployment is needed. Conversely, LSTMs are better suited for datasets with deep temporal structures, such as sales influenced by annual cycles or multi-year trends, despite their higher resource demands.

The optimal strategy often involves blending these tools—using ARIMA for baseline forecasts, XGBoost for variable-rich scenarios, and LSTMs for time-sensitive insights. By tailoring the approach to the business's goals and data profile, retailers can sharpen their forecasting precision, driving smarter decisions and operational excellence.